



Image Classification (Cat/Dog)

The aim of this project is to build and evaluate a robust binary image classification model that can accurately distinguish between cats and dogs using a Convolutional Neural Network (CNN)

```
In [1]: # import the required libraries
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ignore harmless warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [3]: # Importing the data
# Importing the training data
x_train = np.loadtxt('/content/drive/MyDrive/Colab Notebooks/Datasets/Cat and Dog/training_set.csv', delimiter = ',')
y_train = np.loadtxt('/content/drive/MyDrive/Colab Notebooks/Datasets/Cat and Dog/training_set_labels.csv', delimiter = ',')

# Importing the testing data
x_test = np.loadtxt('/content/drive/MyDrive/Colab Notebooks/Datasets/Cat and Dog/test_set.csv', delimiter = ',')
y_test = np.loadtxt('/content/drive/MyDrive/Colab Notebooks/Datasets/Cat and Dog/test_set_labels.csv', delimiter = ',')
```

```
In [4]: x_train
```

```
Out[4]: array([[ 37.,  39.,  25., ...,  58.,  54.,  29.],
   [131., 128., 135., ...,  71.,  96.,  74.],
   [ 80.,  92.,  88., ..., 124., 119.,  99.],
   ...,
   [231., 226., 230., ...,  62.,  65.,  72.],
   [ 61.,  61.,  63., ..., 135., 123., 123.],
   [ 64.,  31.,  12., ...,  61.,  49.,  35.]])
```

```
In [5]: y_train # 0 - Dog and 1 - Cat
```

```
Out[5]: array([0., 0., 0., ..., 1., 1., 1.])
```

```
In [6]: x_test
```

```
Out[6]: array([[118.,  82.,  96., ..., 140.,  79.,  16.],
   [223., 211., 163., ..., 70.,  73.,  78.],
   [ 73.,  67.,  43., ..., 222., 211., 165.],
   ...,
   [249., 245., 242., ..., 73.,  72.,  68.],
   [ 97.,  96., 102., ..., 84.,  78.,  80.],
   [ 94.,  66.,  63., ..., 119.,  96.,  80.]])
```

```
In [7]: # Dimension of the data
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

# (2000, 30000) --> 2000 images of each having 30000 columns, each column is a
(2000, 30000)
(2000,)
(400, 30000)
(400,)
```

```
In [8]: # pixel values
print('Minimum pixel value:', x_train.min())
print('Maximum pixel value:', x_train.max())

# This represents the image can be RGB
```

```
Minimum pixel value: 0.0
Maximum pixel value: 255.0
```

```
In [9]: # reshape the training data
x_train = x_train.reshape(len(x_train), 100, 100, 3)
y_train = y_train.reshape(len(y_train), 1)

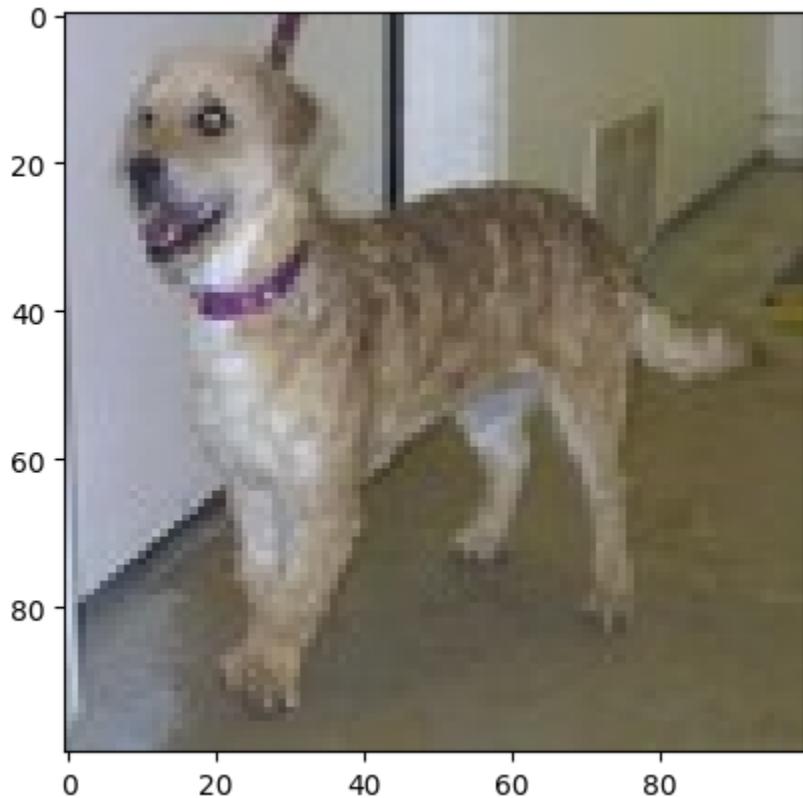
# reshape the test data
x_test = x_test.reshape(len(x_test), 100, 100, 3)
y_test = y_test.reshape(len(y_test), 1)
```

```
In [10]: # data after reshaping
print('Shape of the data after reshaping:')
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
```

```
Shape of the data after reshaping:
(2000, 100, 100, 3)
(2000, 1)
(400, 100, 100, 3)
(400, 1)
```

```
In [11]: # Data scaling
x_train = x_train/255.0
x_test = x_test/255.0
```

```
In [12]: # printing the first image  
plt.imshow(x_train[355, :])  
plt.show()  
  
# print the target for the first image  
print('The target label is:', y_train[355])
```



The target label is: [0.]

CNN Model Building

```
In [13]: # import deep learning libraries  
from sklearn.metrics import confusion_matrix  
from sklearn.model_selection import train_test_split  
import keras  
from keras.applications import MobileNetV2  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
from keras.models import Model  
from keras.layers import GlobalAveragePooling2D, Dense, Dropout
```

```
In [14]: # Creating a validation split  
x_train_new, x_val, y_train_new, y_val = train_test_split(x_train, y_train, te
```

```
In [15]: base_model = MobileNetV2(  
    weights='imagenet',  
    include_top=False,  
    input_shape=(100, 100, 3)
```

```
)  
base_model.trainable = False
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5
9406464/9406464 ————— 0s 0us/step

```
In [16]: #Adding a clean classifier head  
x = base_model.output  
x = GlobalAveragePooling2D()(x)  
x = Dense(128, activation='relu')(x)  
x = Dropout(0.5)(x)  
output = Dense(1, activation='sigmoid')(x)  
  
model = Model(inputs=base_model.input, outputs=output)
```

```
In [17]: # Compiling the model  
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
In [18]: # Adding data augmentation  
datagen = ImageDataGenerator(  
    rotation_range=20,  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    zoom_range=0.2,  
    horizontal_flip=True  
)  
  
datagen.fit(x_train_new)
```

```
In [19]: # training the model with the data  
history = model.fit(datagen.flow(x_train_new, y_train_new, batch_size=32), epochs=10)
```


```
Epoch 1/10
57/57 ━━━━━━━━━━ 44s 471ms/step - accuracy: 0.7416 - loss: 0.5790 - val_accuracy: 0.8750 - val_loss: 0.2534
Epoch 2/10
57/57 ━━━━━━━━━━ 5s 81ms/step - accuracy: 0.8562 - loss: 0.3001 - val_accuracy: 0.9000 - val_loss: 0.2434
Epoch 3/10
57/57 ━━━━━━━━━━ 6s 101ms/step - accuracy: 0.8903 - loss: 0.2526 - val_accuracy: 0.9000 - val_loss: 0.2434
Epoch 4/10
57/57 ━━━━━━━━━━ 4s 75ms/step - accuracy: 0.8876 - loss: 0.2445 - val_accuracy: 0.9050 - val_loss: 0.2683
Epoch 5/10
57/57 ━━━━━━━━━━ 4s 75ms/step - accuracy: 0.9251 - loss: 0.1990 - val_accuracy: 0.9000 - val_loss: 0.2521
Epoch 6/10
57/57 ━━━━━━━━━━ 6s 99ms/step - accuracy: 0.8974 - loss: 0.2279 - val_accuracy: 0.9100 - val_loss: 0.2523
Epoch 7/10
57/57 ━━━━━━━━━━ 4s 74ms/step - accuracy: 0.9081 - loss: 0.2210 - val_accuracy: 0.9200 - val_loss: 0.2489
Epoch 8/10
57/57 ━━━━━━━━━━ 5s 82ms/step - accuracy: 0.9040 - loss: 0.1974 - val_accuracy: 0.8950 - val_loss: 0.2767
Epoch 9/10
57/57 ━━━━━━━━━━ 5s 90ms/step - accuracy: 0.9023 - loss: 0.2042 - val_accuracy: 0.8950 - val_loss: 0.2667
Epoch 10/10
57/57 ━━━━━━━━━━ 4s 75ms/step - accuracy: 0.9147 - loss: 0.1919 - val_accuracy: 0.9250 - val_loss: 0.2677
```

```
In [20]: # evaluating the model
model.evaluate(x_test, y_test)
```

```
13/13 ━━━━━━━━━━ 10s 817ms/step - accuracy: 0.9119 - loss: 0.2116
```

```
Out[20]: [0.21065618097782135, 0.9175000190734863]
```


```
In [21]: # Training vs validation accuracy
plt.plot(history.history['accuracy'], label='Train')
plt.plot(history.history['val_accuracy'], label='Validation')
plt.legend()
plt.show()
```



```
In [22]: # Confusion matrix
y_prob = model.predict(x_test)
y_pred = (y_prob >= 0.5).astype(int).reshape(-1)
cm = confusion_matrix(y_test.reshape(-1), y_pred)

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

13/13 ————— 8s 350ms/step



```
In [23]: # training data distribution of target
unique, counts = np.unique(y_train, return_counts=True)
print(unique, counts)
```

```
[0. 1.] [1000 1000]
```

```
In [24]: len(x_test)
```

```
Out[24]: 400
```


```
In [26]: # predictions
rdm_img = random.randint(0, len(x_test))
plt.imshow(x_test[rdm_img, :])
plt.show()

# actual condition
if y_test[rdm_img, :] == 0:
    actual = 'Dog'
else:
    actual = 'Cat'

# prediction for the given image
y_pred = model.predict(x_test[rdm_img, :].reshape(1, 100, 100, 3))
print('Prediction Probability:', y_pred)

if y_pred >= 0.5:
    pred = 'Cat'
```

```
else:  
    pred = 'Dog'  
print('The actual image is:', actual)  
print('The model has predicted as:', pred)  
print('The image value', rdm_img)
```



```
1/1 ————— 0s 33ms/step  
Prediction Probability: [[0.9878637]]  
The actual image is: Cat  
The model has predicted as: Cat  
The image value 248
```

Final Conclusion & Limitations

Final Conclusion

I have used MobileNetV2 for binary image classification (cats vs dogs), the model achieved ~92% test accuracy and showed stable training-validation behavior. The results confirm that leveraging pretrained CNNs can deliver strong performance even with a relatively small dataset, making this approach practical and efficient for real-world image classification tasks.

Limitations

- The dataset is small and controlled, so performance may drop on real-world images with different lighting, backgrounds, or breeds.
- Only the classifier head was trained; the base model was not fine-tuned, which may limit peak performance.
- This setup handles binary classification only and would need architectural changes for multi-class problems.

In [25]: