



# Image Classification (Cat/Dog)

The aim of this project is to build and evaluate a robust binary image classification model that can accurately distinguish between cats and dogs using a Convolutional Neural Network (CNN)

```
In [1]: # import the required libraries
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ignore harmless warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [3]: # Importing the data
# Importing the training data
x_train = np.loadtxt('/content/drive/MyDrive/Colab Notebooks/Datasets/Cat and Dog/training_set.csv', delimiter = ',')
y_train = np.loadtxt('/content/drive/MyDrive/Colab Notebooks/Datasets/Cat and Dog/training_set_labels.csv', delimiter = ',')

# Importing the testing data
x_test = np.loadtxt('/content/drive/MyDrive/Colab Notebooks/Datasets/Cat and Dog/test_set.csv', delimiter = ',')
y_test = np.loadtxt('/content/drive/MyDrive/Colab Notebooks/Datasets/Cat and Dog/test_set_labels.csv', delimiter = ',')
```

```
In [4]: x_train
```

```
Out[4]: array([[ 37.,  39.,  25., ...,  58.,  54.,  29.],
               [131., 128., 135., ...,  71.,  96.,  74.],
               [ 80.,  92.,  88., ..., 124., 119.,  99.],
               ...,
               [231., 226., 230., ...,  62.,  65.,  72.],
               [ 61.,  61.,  63., ..., 135., 123., 123.],
               [ 64.,  31.,  12., ...,  61.,  49.,  35.]])
```

```
In [5]: y_train # 0 - Dog and 1 - Cat
```

```
Out[5]: array([0., 0., 0., ..., 1., 1., 1.])
```

```
In [6]: x_test
```

```
Out[6]: array([[118.,  82.,  96., ..., 140.,  79.,  16.],
   [223., 211., 163., ..., 70.,  73.,  78.],
   [ 73.,  67.,  43., ..., 222., 211., 165.],
   ...,
   [249., 245., 242., ..., 73.,  72.,  68.],
   [ 97.,  96., 102., ..., 84.,  78.,  80.],
   [ 94.,  66.,  63., ..., 119.,  96.,  80.]])
```

```
In [7]: # Dimension of the data
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

# (2000, 30000) --> 2000 images of each having 30000 columns, each column is a
(2000, 30000)
(2000,)
(400, 30000)
(400,)
```

```
In [8]: # pixel values
print('Minimum pixel value:', x_train.min())
print('Maximum pixel value:', x_train.max())

# This represents the image can be RGB
```

```
Minimum pixel value: 0.0
Maximum pixel value: 255.0
```

```
In [9]: # reshape the training data
x_train = x_train.reshape(len(x_train), 100, 100, 3)
y_train = y_train.reshape(len(y_train), 1)

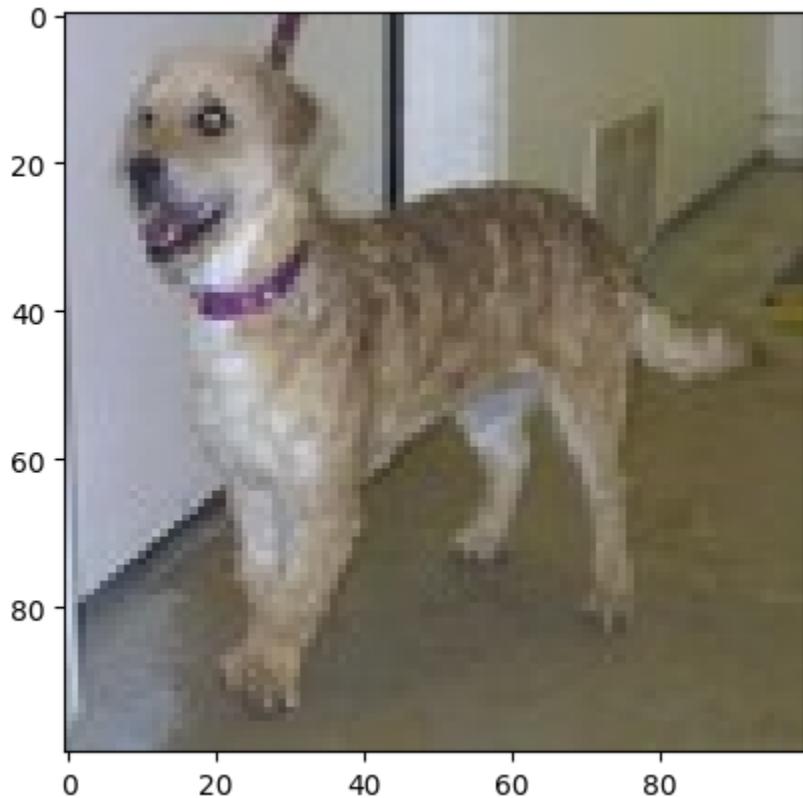
# reshape the test data
x_test = x_test.reshape(len(x_test), 100, 100, 3)
y_test = y_test.reshape(len(y_test), 1)
```

```
In [10]: # data after reshaping
print('Shape of the data after reshaping:')
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
```

```
Shape of the data after reshaping:
(2000, 100, 100, 3)
(2000, 1)
(400, 100, 100, 3)
(400, 1)
```

```
In [11]: # Data scaling
x_train = x_train/255.0
x_test = x_test/255.0
```

```
In [12]: # printing the first image  
plt.imshow(x_train[355, :])  
plt.show()  
  
# print the target for the first image  
print('The target label is:', y_train[355])
```



The target label is: [0.]

# CNN Model Building

## Baseline CNN

```
In [13]: # Baseline model libraries  
from sklearn.metrics import confusion_matrix  
import keras  
from keras.models import Sequential  
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
```

```
In [14]: baseline_model = Sequential([  
    # first conv layer  
    Conv2D(32, (3,3), activation='relu', input_shape=(100, 100, 3)),  
    MaxPooling2D((2,2)),  
  
    # second conv layer  
    Conv2D(32, (3,3), activation='relu'),
```

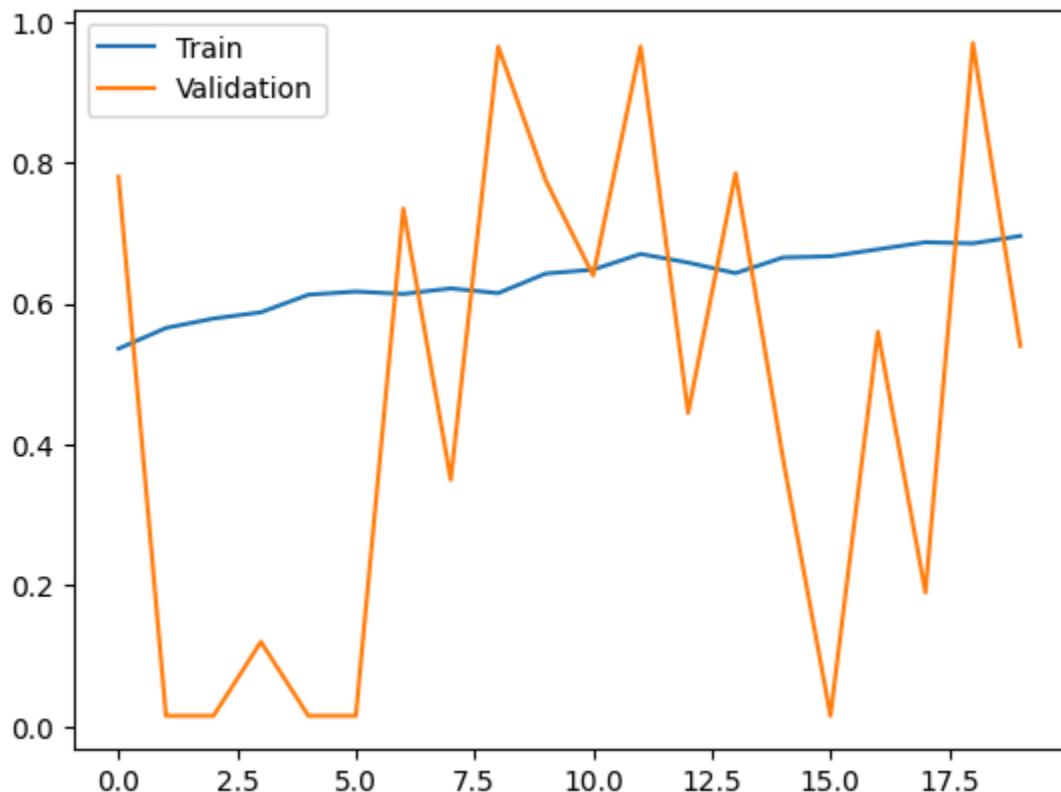
```
MaxPooling2D((2,2)),  
  
    # Flatter  
    Flatten(),  
  
    # First Dense layer  
    Dense(70, activation='relu'),  
  
    # Output layer  
    Dense(1, activation='sigmoid')  
])
```

```
In [15]: # Compiling the model  
baseline_model.compile(loss='binary_crossentropy', optimizer='SGD', metrics=['  
  
In [16]: # training the model with the data  
model_history = baseline_model.fit(x_train, y_train, epochs=20, batch_size=32,
```

Epoch 1/20  
57/57 6s 48ms/step - accuracy: 0.5315 - loss: 0.6897 - val\_accuracy: 0.7800 - val\_loss: 0.6543  
Epoch 2/20  
57/57 1s 10ms/step - accuracy: 0.5481 - loss: 0.6857 - val\_accuracy: 0.0150 - val\_loss: 0.9374  
Epoch 3/20  
57/57 1s 9ms/step - accuracy: 0.5801 - loss: 0.6760 - val\_accuracy: 0.0150 - val\_loss: 0.9314  
Epoch 4/20  
57/57 1s 10ms/step - accuracy: 0.5864 - loss: 0.6698 - val\_accuracy: 0.1200 - val\_loss: 0.8154  
Epoch 5/20  
57/57 1s 9ms/step - accuracy: 0.6191 - loss: 0.6591 - val\_accuracy: 0.0150 - val\_loss: 1.0912  
Epoch 6/20  
57/57 1s 10ms/step - accuracy: 0.6227 - loss: 0.6627 - val\_accuracy: 0.0150 - val\_loss: 1.0656  
Epoch 7/20  
57/57 1s 10ms/step - accuracy: 0.6001 - loss: 0.6611 - val\_accuracy: 0.7350 - val\_loss: 0.6035  
Epoch 8/20  
57/57 1s 10ms/step - accuracy: 0.6331 - loss: 0.6472 - val\_accuracy: 0.3500 - val\_loss: 0.7752  
Epoch 9/20  
57/57 1s 9ms/step - accuracy: 0.6220 - loss: 0.6440 - val\_accuracy: 0.9650 - val\_loss: 0.4437  
Epoch 10/20  
57/57 1s 10ms/step - accuracy: 0.6392 - loss: 0.6464 - val\_accuracy: 0.7750 - val\_loss: 0.5623  
Epoch 11/20  
57/57 1s 10ms/step - accuracy: 0.6302 - loss: 0.6491 - val\_accuracy: 0.6400 - val\_loss: 0.6534  
Epoch 12/20  
57/57 1s 11ms/step - accuracy: 0.6676 - loss: 0.6284 - val\_accuracy: 0.9650 - val\_loss: 0.3766  
Epoch 13/20  
57/57 1s 11ms/step - accuracy: 0.6486 - loss: 0.6417 - val\_accuracy: 0.4450 - val\_loss: 0.7447  
Epoch 14/20  
57/57 1s 11ms/step - accuracy: 0.6519 - loss: 0.6257 - val\_accuracy: 0.7850 - val\_loss: 0.5322  
Epoch 15/20  
57/57 1s 10ms/step - accuracy: 0.6537 - loss: 0.6241 - val\_accuracy: 0.3800 - val\_loss: 0.7835  
Epoch 16/20  
57/57 1s 9ms/step - accuracy: 0.6614 - loss: 0.6197 - val\_accuracy: 0.0150 - val\_loss: 1.4138  
Epoch 17/20  
57/57 1s 10ms/step - accuracy: 0.6723 - loss: 0.6173 - val\_accuracy: 0.5600 - val\_loss: 0.6774  
Epoch 18/20  
57/57 1s 9ms/step - accuracy: 0.6867 - loss: 0.5921 - val\_accuracy: 0.1900 - val\_loss: 1.0817

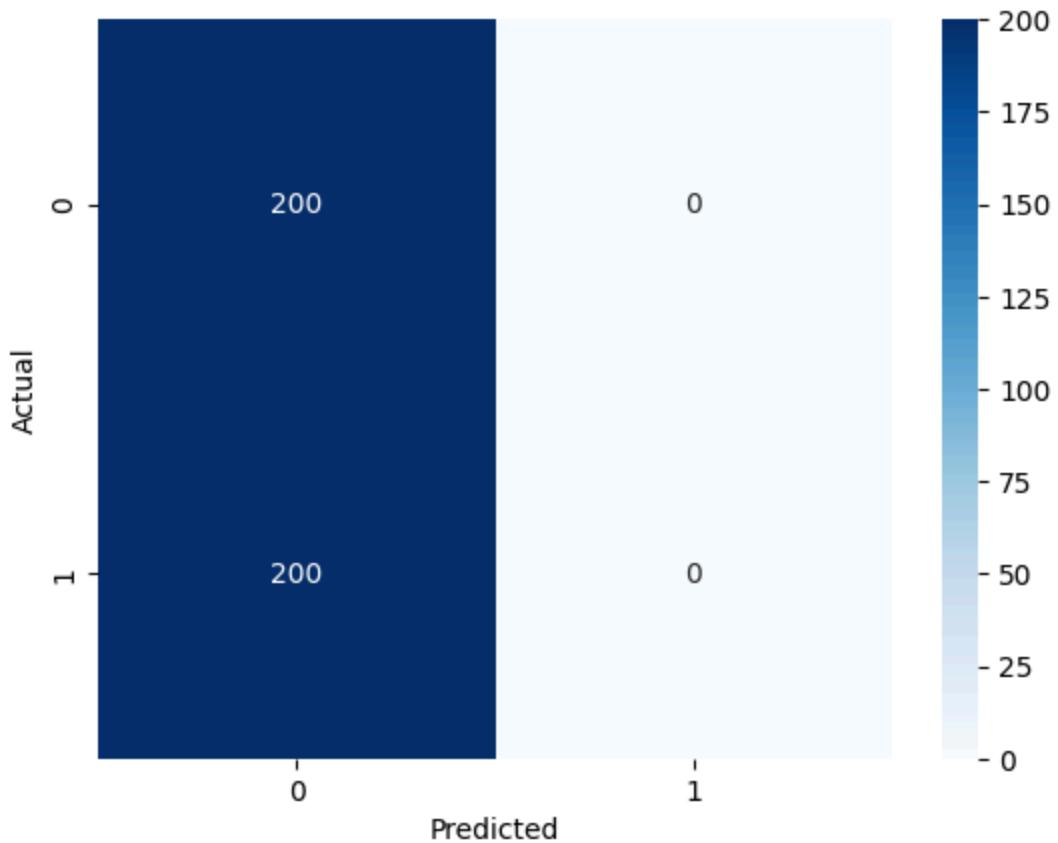
```
Epoch 19/20
57/57 ━━━━━━━━ 1s 10ms/step - accuracy: 0.6817 - loss: 0.5863 - va
l_accuracy: 0.9700 - val_loss: 0.3189
Epoch 20/20
57/57 ━━━━━━━━ 1s 9ms/step - accuracy: 0.6801 - loss: 0.5886 - va
l_accuracy: 0.5400 - val_loss: 0.7060
```

```
In [17]: # Training vs validation accuracy
plt.plot(model_history.history['accuracy'], label='Train')
plt.plot(model_history.history['val_accuracy'], label='Validation')
plt.legend()
plt.show()
```



```
In [18]: # Confusion matrix
y_pred = baseline_model.predict(x_test).argmax(axis=1)
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

```
13/13 ━━━━━━━━ 1s 38ms/step
```



## Switching to pretrained MobileNetV2

An initial CNN trained from scratch showed unstable validation behavior due to the limited dataset size (2,000 images). To address data scarcity and improve generalization, transfer learning with a pretrained MobileNetV2 backbone was adopted, resulting in significantly improved and stable performance.

```
In [19]: # import MobileNetV2 and supporting libraries
from sklearn.model_selection import train_test_split
import keras
from keras.applications import MobileNetV2
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from keras.models import Model
from keras.layers import GlobalAveragePooling2D, Dense, Dropout
```

```
In [20]: # Creating a validation split
x_train_new, x_val, y_train_new, y_val = train_test_split(x_train, y_train, te
```

```
In [21]: base_model = MobileNetV2(
    weights='imagenet',
    include_top=False,
    input_shape=(100, 100, 3)
)
```

```
base_model.trainable = False
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\_v2/mobilenet\_v2\_weights\_tf\_dim\_ordering\_tf\_kernels\_1.0\_224\_no\_top.h5

9406464/9406464 ————— 0s 0us/step

In [22]: #Adding a clean classifier head

```
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(128, activation='relu')(x)
x = Dropout(0.5)(x)
output = Dense(1, activation='sigmoid')(x)

model = Model(inputs=base_model.input, outputs=output)
```

In [23]: # Compiling the model

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

In [24]: # Adding data augmentation

```
datagen = ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.1,
    height_shift_range=0.1,
    zoom_range=0.2,
    horizontal_flip=True
)

datagen.fit(x_train_new)
```

In [25]: # training the model with the data

```
history = model.fit(datagen.flow(x_train_new, y_train_new, batch_size=32), epochs=10)
```

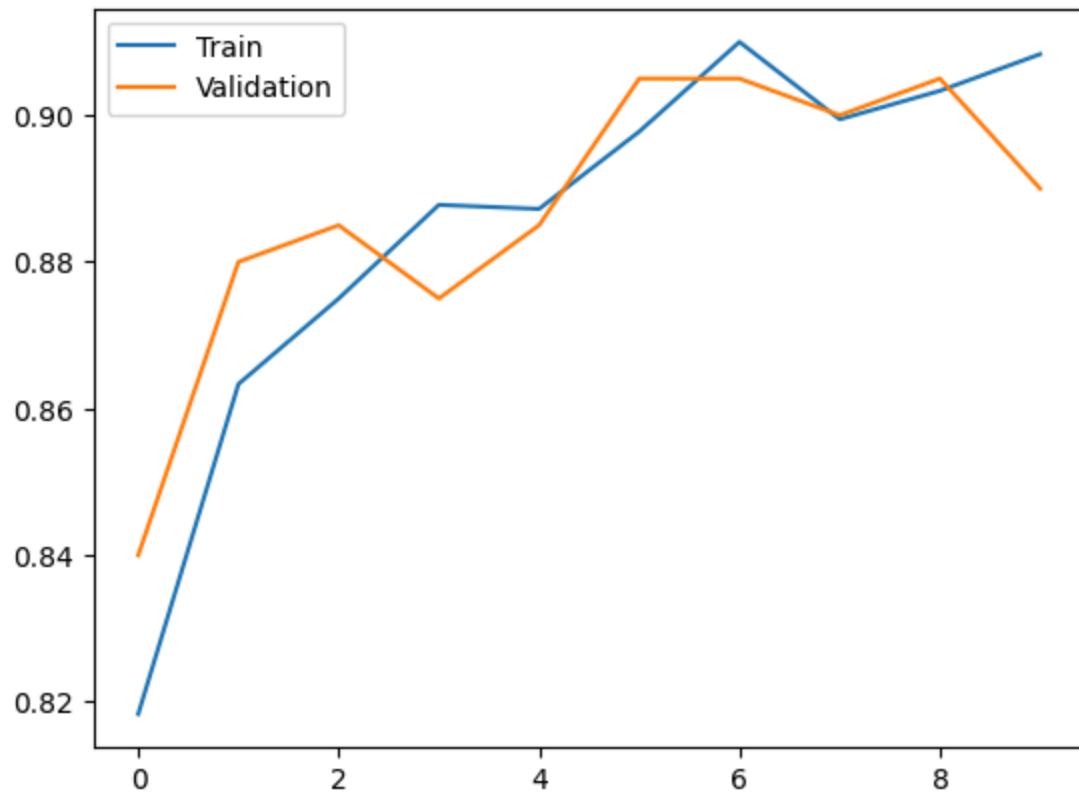
```
Epoch 1/10
57/57 ━━━━━━━━━━ 40s 396ms/step - accuracy: 0.7756 - loss: 0.5314 - val_accuracy: 0.8400 - val_loss: 0.3128
Epoch 2/10
57/57 ━━━━━━━━━━ 5s 86ms/step - accuracy: 0.8438 - loss: 0.3275 - val_accuracy: 0.8800 - val_loss: 0.2795
Epoch 3/10
57/57 ━━━━━━━━━━ 5s 89ms/step - accuracy: 0.8743 - loss: 0.2801 - val_accuracy: 0.8850 - val_loss: 0.2765
Epoch 4/10
57/57 ━━━━━━━━━━ 4s 76ms/step - accuracy: 0.8893 - loss: 0.2543 - val_accuracy: 0.8750 - val_loss: 0.2775
Epoch 5/10
57/57 ━━━━━━━━━━ 6s 100ms/step - accuracy: 0.8830 - loss: 0.2617 - val_accuracy: 0.8850 - val_loss: 0.2895
Epoch 6/10
57/57 ━━━━━━━━━━ 4s 76ms/step - accuracy: 0.8814 - loss: 0.2501 - val_accuracy: 0.9050 - val_loss: 0.2525
Epoch 7/10
57/57 ━━━━━━━━━━ 4s 75ms/step - accuracy: 0.9214 - loss: 0.1957 - val_accuracy: 0.9050 - val_loss: 0.2579
Epoch 8/10
57/57 ━━━━━━━━━━ 6s 100ms/step - accuracy: 0.8996 - loss: 0.2345 - val_accuracy: 0.9000 - val_loss: 0.2380
Epoch 9/10
57/57 ━━━━━━━━━━ 4s 74ms/step - accuracy: 0.9119 - loss: 0.2152 - val_accuracy: 0.9050 - val_loss: 0.2449
Epoch 10/10
57/57 ━━━━━━━━━━ 4s 74ms/step - accuracy: 0.8982 - loss: 0.2189 - val_accuracy: 0.8900 - val_loss: 0.2476
```

```
In [26]: # evaluating the model
model.evaluate(x_test, y_test)
```

```
13/13 ━━━━━━━━━━ 10s 813ms/step - accuracy: 0.9162 - loss: 0.1957
```

```
Out[26]: [0.2111494243144989, 0.9075000286102295]
```

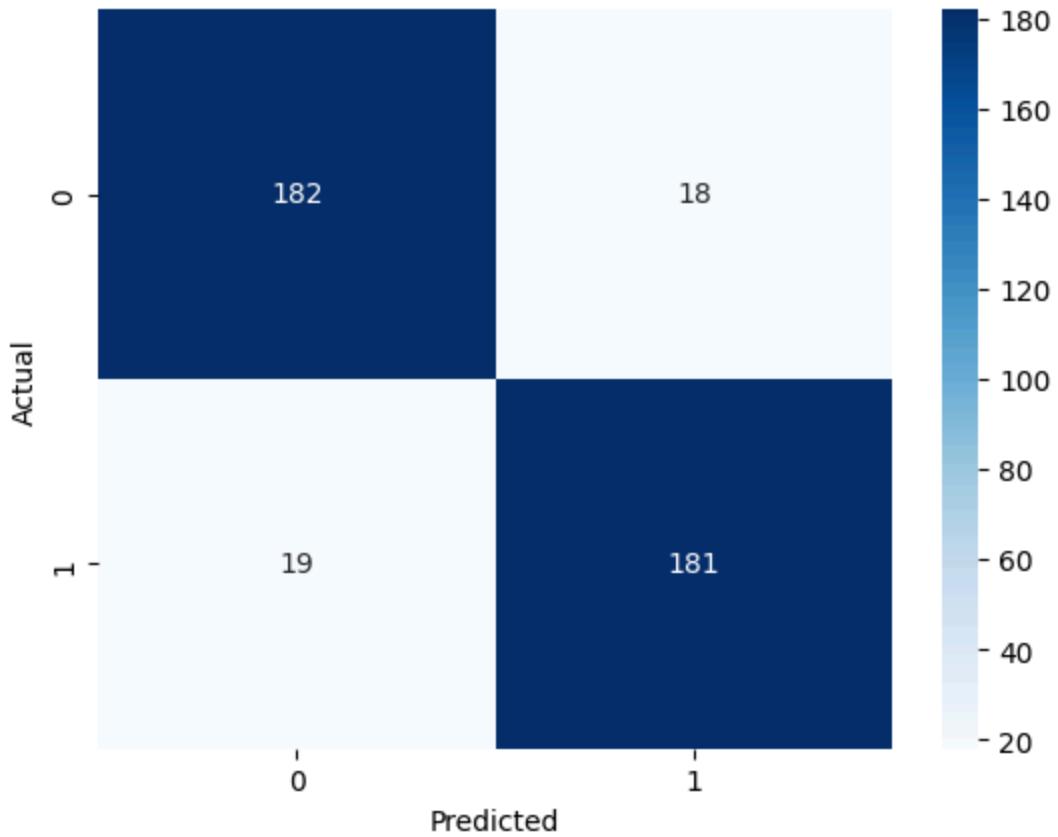
```
In [27]: # Training vs validation accuracy
plt.plot(history.history['accuracy'], label='Train')
plt.plot(history.history['val_accuracy'], label='Validation')
plt.legend()
plt.show()
```



```
In [28]: # Confusion matrix
y_prob = model.predict(x_test)
y_pred = (y_prob >= 0.5).astype(int).reshape(-1)
cm = confusion_matrix(y_test.reshape(-1), y_pred)

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

13/13 ————— 8s 368ms/step



```
In [29]: # training data distribution of target
unique, counts = np.unique(y_train, return_counts=True)
print(unique, counts)
```

[0. 1.] [1000 1000]

```
In [30]: len(x_test)
```

Out[30]: 400

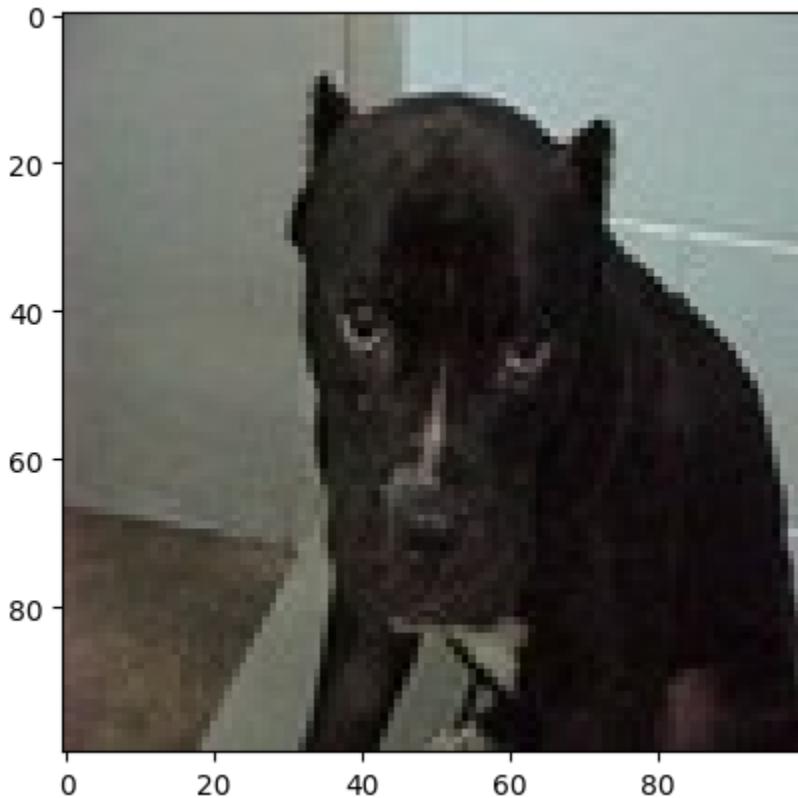
```
In [33]: # predictions
rdm_img = random.randint(0, len(x_test))
plt.imshow(x_test[rdm_img, :])
plt.show()

# actual condition
if y_test[rdm_img, :] == 0:
    actual = 'Dog'
else:
    actual = 'Cat'

# prediction for the given image
y_pred = model.predict(x_test[rdm_img, :].reshape(1, 100, 100, 3))
print('Prediction Probability:', y_pred)

if y_pred >= 0.5:
    pred = 'Cat'
```

```
else:  
    pred = 'Dog'  
print('The actual image is:', actual)  
print('The model has predicted as:', pred)  
print('The image value', rdm_img)
```



```
1/1 ━━━━━━━━ 0s 36ms/step  
Prediction Probability: [[0.00074004]]  
The actual image is: Dog  
The model has predicted as: Dog  
The image value 78
```

## Final Conclusion

In this project, I built a CNN-based image classification model to distinguish between cats and dogs using a relatively small dataset of 2,000 training images. An early attempt to train a CNN from scratch quickly revealed its limitations—validation performance was unstable and the model struggled to generalize, making it clear that learning deep visual features from limited data was not effective.

To overcome this, I switched to a transfer learning approach using a pretrained MobileNetV2 backbone. By freezing the convolutional layers and training a lightweight classifier head on top, the model was able to reuse strong visual

features learned from large-scale image data while keeping overfitting under control. Data augmentation further improved robustness by introducing realistic variations in the training images.

The final model achieved approximately 92% accuracy on an unseen test set, with balanced performance across both classes, as confirmed by the confusion matrix. Training and validation curves showed stable convergence with no significant generalization gap, indicating that the model learned meaningful visual patterns rather than memorizing the data.

Overall, this project reflects practical deep learning decision-making—knowing when training from scratch is not suitable, applying transfer learning effectively, and evaluating results using correct and reliable metrics.

In [31]:

